

Energy-Based Generative Adversarial Networks

Wenhai Yang

School of Mathematical Sciences
Peking University

March 9, 2017

Outline

1 Introduction

- Energy-Based Model
- Generative Adversarial Networks

2 Energy-Based Generative Adversarial Networks

3 EBGAN Model

4 Experiments Results

5 Reference

Energy-Based Model

- Try to learn a function f , which maps each point of an input space to a single scalar.

$$f : X \rightarrow R \tag{1}$$

where the scalar is called “energy”.

- Ideal function f : the desired configuration gets assigned low energies while the incorrect ones are given high energies.
- E.g:
 - Supervised learning: low energy corresponds to correct label.
 - Unsupervised learning: low energy corresponds to the data manifold.

Generative Adversarial Networks

GAN are combined with the following parts:

- Discriminator
- Generator
- etc

Energy-Based Generative Adversarial Networks

- Discriminator is viewed as an energy function, without explicit probabilistic interpretation.
- Discriminator is trained to assign low energy values to regions of high data density.
- Generator is viewed as a trainable parameterized function that produces samples in regions of the space to which the generator assigns low energy.

EBGAN Model

- Notation:

p_{data} : the underlying probability density of the distribution that produces the dataset.

G : the generator which is trained to produce a sample $G(z)$

p_z : A known distribution to provide a random vector z

D : the discriminator which takes either real or generated samples and estimates the energy value $E \in R$

- Objectie functional:

Given a positive margin m , a data sample x and a generated sample $G(z)$, defining the discriminator loss \mathcal{L}_D and the generator loss \mathcal{L}_G :

$$\mathcal{L}_D(x, z) = D(x) + [m - D(G(z))]^+ \quad (2)$$

$$\mathcal{L}_G(z) = D(G(z)) \quad (3)$$

EBGAN Model

- Define:

$$V(G, D) = \int_{x,z} \mathcal{L}_D(x, z)p_{data}(x)p_z(z)dxdz \quad (4)$$

$$U(G, D) = \int_z \mathcal{L}_G(z)p_z(z)dz \quad (5)$$

- Aim: train the D to minimize V and train the G to minimize U
- Nash equilibrium of the system is a pair (G^*, D^*) that satisfies:

$$V(G^*, D^*) \leq V(G^*, D) \quad \forall D \quad (6)$$

$$U(G^*, D^*) \leq U(G, D^*) \quad \forall G \quad (7)$$

EBGAN Model

Theorem

If (D^*, G^*) is a Nash equilibrium of the system, then $p_G \stackrel{a.s}{=} p_{data}$ and $V(D^*, G^*) = m$

Theorem

Nash equilibrium of this system exists and is characterized by:

- (a) $p_G \stackrel{a.s}{=} p_{data}$
- (b) there exists a constant $\gamma \in [0, m]$ such that $D^*(x) \stackrel{a.s}{=} \gamma$

Auto-Encoders

The discriminator D is structured as an auto-encoder:

$$D(x) = \|Dec(Enc(x)) - x\| \quad (8)$$

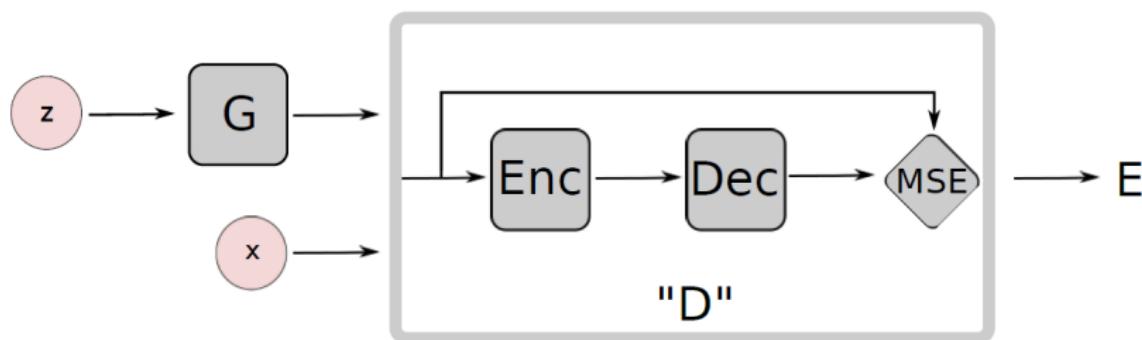


Figure 1: EBGAN architecture with an auto-encoder discriminator.

Auto-Encoders

Why Auto-Encoders?

- Offer a diverse targets for the discriminator
- Reconstruction loss will likely produce very different gradient directions within the minibatch, allowing for larger minibatch size without loss of efficiency.
- Given regularization, auto-encoders have the ability to learn an energy manifold without supervision or negative examples.

Problems:

- Auto-encoders may learn little more than an identity function, which means attributing low energy to the whole space.

Solution:

- Regularization: to give higher energy to points outside of the data manifold.

Auto-Encoders

Repelling Regularizer:

- Aim: to keep the model from producing samples that are clustered in one or a few modes of p_{data}
- Pulling-away(PT) effect:

$$f_{PT}(S) = \frac{1}{N(N-1)} \sum_i \sum_{j \neq i} \left(\frac{S_i^T S_j}{\|S_i\| \|S_j\|} \right)^2 \quad (9)$$

where $S \in R^{s \times N}$ denotes a batch of sample representations taken from the encoder output layer.

- PT term is used in the generator loss but not discriminator loss

Exhaustive Grid Search On MNIST

Table 1: Grid search specs

Settings	Description	EBGANs	GANs
nLayerG	number of layers in G	[2, 3, 4, 5]	[2, 3, 4, 5]
nLayerD	number of layers in D	[2, 3, 4, 5]	[2, 3, 4, 5]
sizeG	number of neurons in G	[400, 800, 1600, 3200]	[400, 800, 1600, 3200]
sizeD	number of neurons in D	[128, 256, 512, 1024]	[128, 256, 512, 1024]
dropoutD	if to use dropout in D	[true, false]	[true, false]
optimD	to use Adam or SGD for D	adam	[adam, sgd]
optimG	to use Adam or SGD for G	adam	[adam, sgd]
lr	learning rate	0.001	[0.01, 0.001, 0.0001]
#experiments:	-	512	6144

Exhaustive Grid Search On MNIST

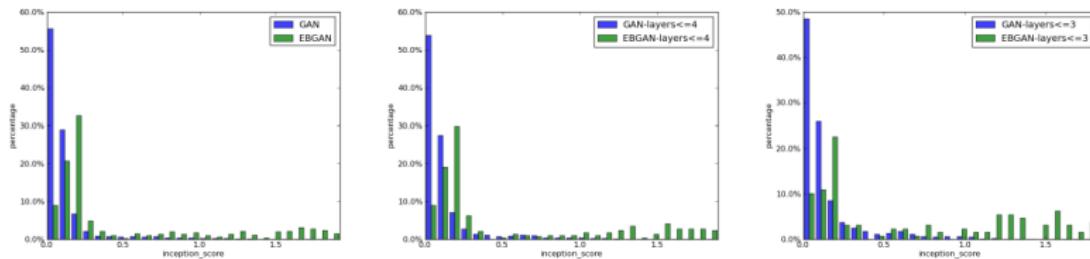


Figure 2: **(Zooming in on pdf file is recommended.)** Histogram of the inception scores from the grid search. The x-axis carries the inception score I and y-axis informs the portion of the models (in percentage) falling into certain bins. Left (a): general comparison of EBGANs against GANs; Middle (b): EBGANs and GANs both constrained by $n\text{Layer} [\text{GD}] \leq 4$; Right (c): EBGANs and GANs both constrained by $n\text{Layer} [\text{GD}] \leq 3$.

Exhaustive Grid Search On MNITST

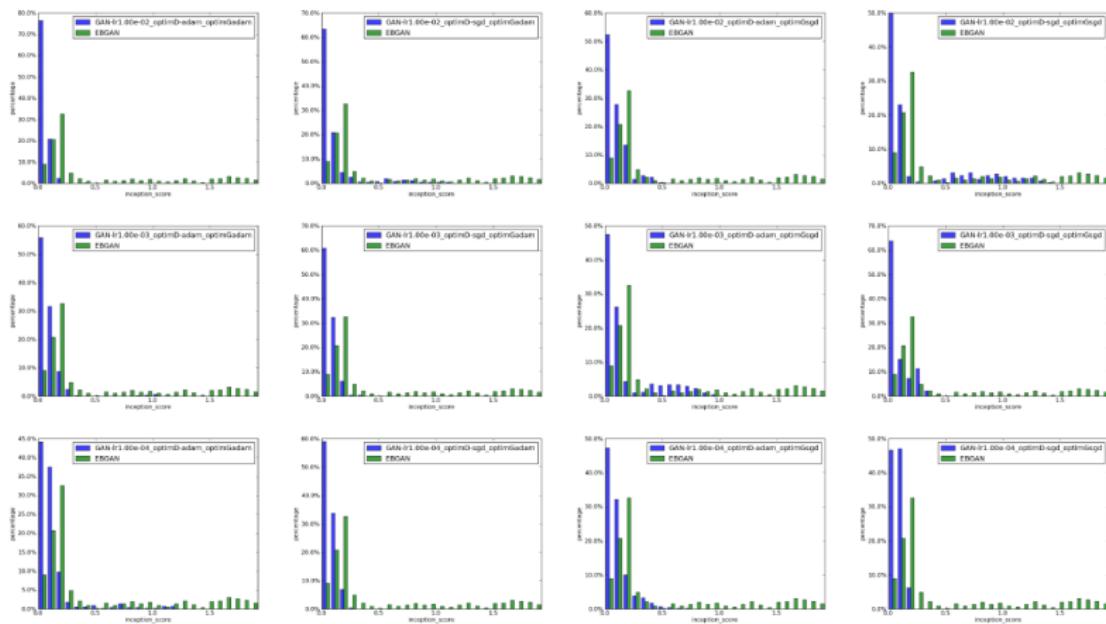


Figure 3: **(Zooming in on pdf file is recommended.)** Histogram of the inception scores grouped by different optimization combinations, drawn from optimD, optimG and lr (See text).

Exhaustive Grid Search On MNIST



Figure 1: Best GAN model

Exhaustive Grid Search On MNIST



Figure 2: Best EBGAN model

Exhaustive Grid Search On MNIST



Figure 3: Best EBGAN-PT model

Semi-Supervised Learning On MNIST

Positioning a bottom-layer-cost LN into an EBGAN framework.

Table 2: The comparison of LN bottom-layer-cost model and its EBGAN extension on PI-MNIST semi-supervised task. Note the results are error rate (in %) and they were averaged over 15 different random seeds.

model	100	200	1000
LN bottom-layer-cost, reported in Pezeshki et al. (2015)	1.69 ± 0.18	-	1.05 ± 0.02
LN bottom-layer-cost, reported in Rasmus et al. (2015)	1.09 ± 0.32	-	0.90 ± 0.05
LN bottom-layer-cost, reproduced in this work (see appendix D)	1.36 ± 0.21	1.24 ± 0.09	1.04 ± 0.06
LN bottom-layer-cost within EBGAN framework	1.04 ± 0.12	0.99 ± 0.12	0.89 ± 0.04
Relative percentage improvement	23.5%	20.2%	14.4%

LSUN & CELEBA

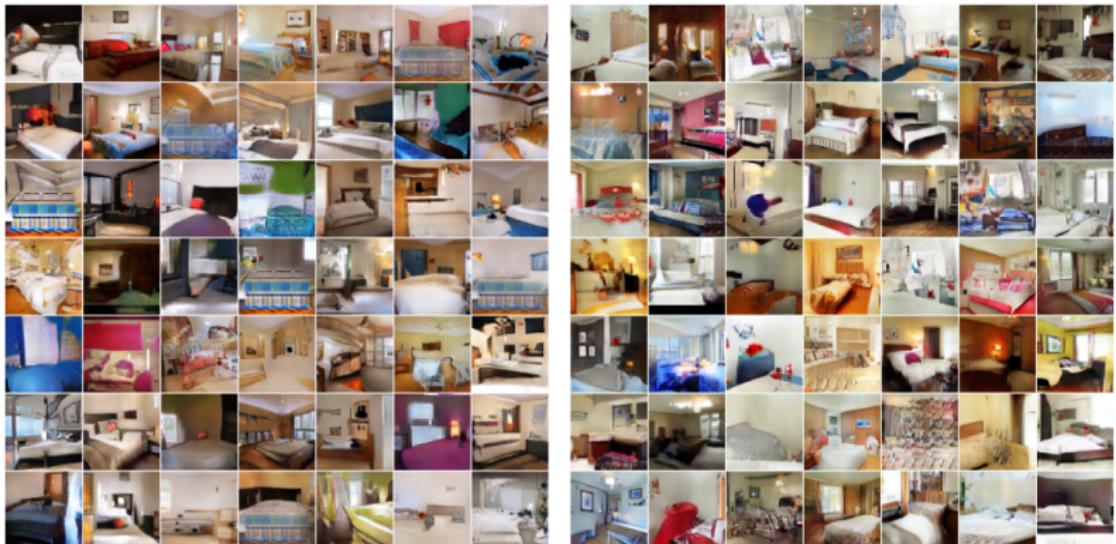


Figure 5: Generation from LSUN bedroom full-images. Left(a): DCGAN generation. Right(b): EBGAN-PT generation.

LSUN & CELEBA



Figure 6: Generation from CelebA face dataset. Left(a): DCGAN generation. Right(b): EBGAN-PT generation.

Imagenet

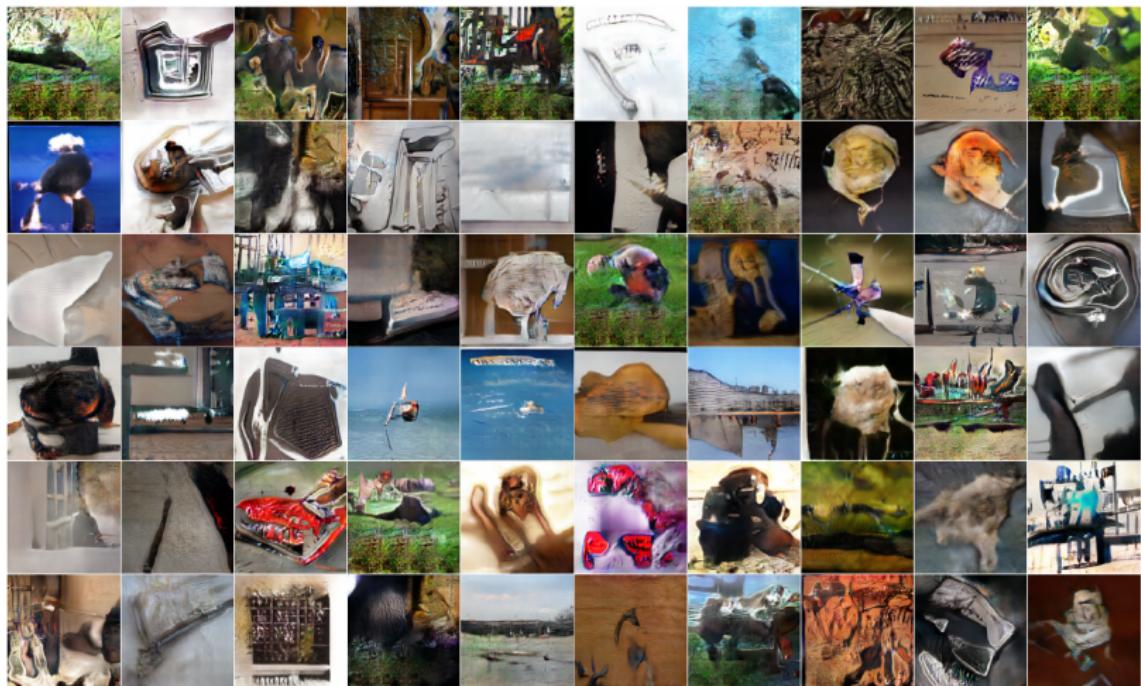


Figure 7: ImageNet 128×128 generations using an EBGAN-PT.

Imagenet



Figure 8: ImageNet 256×256 generations using an EBGAN-PT.

Reference

Reference:

- [1] J ZhaoM MathieuY Lecun; Energy-based Generative Adversarial Network